

MOHAMMED YAHYA MOHAMMED QAYYUM QURESHI

Mumbai, India • +91 9326780873 • yahyaqureshi677@gmail.com

[LinkedIn](#) • [Portfolio](#) • [GitHub](#) • [Kaggle](#)

PROFESSIONAL SUMMARY

Data Scientist and AI/ML developer specializing in multi-agent systems and Retrieval-Augmented Generation (RAG), with a 9.28 CGPA and hands-on experience shipping full-stack ML applications from data pipeline to deployment. Built a production-style multi-agent trading intelligence platform using persistent knowledge graphs and vector retrieval, and validated a production-data ML pipeline at FlyRank.ai using client-grouped evaluation and leakage audits.

SKILLS & ABILITIES

- AI & Machine Learning:** NumPy, Pandas, Matplotlib, Scikit-Learn, Multi-Agent Systems, Retrieval-Augmented Generation (RAG), Knowledge Graphs (LightRAG), Predictive Modeling, and Data Mining (WEKA).
- Data Analytics:** Statistical Analysis (Inferential & Descriptive), Exploratory Data Analysis (EDA), Data Visualization, and Business Intelligence.
- Programming Languages:** Python (Advanced), C, R, and T-SQL.
- Web Frameworks & Tools:** FastAPI, Flask, Streamlit, HTML/CSS, JavaScript.
- Databases:** Hadoop, SQL Server (SSMS), PostgreSQL, MySQL, MongoDB (NoSQL), and Supabase.
- Developer Tools:** Git, Vercel, VectorDB, and MS Excel (Advanced).

PROJECTS

- Prediction Market Trader Advisor** — Lead Developer, AI Multi-Agent Platform Mar–Apr 2026
- Architected a modular multi-agent AI system with a three-stage execution pipeline (Intent Classifier, Query Router, Intelligence Agent) to automate high-conviction trading recommendations for Polymarket and Kalshi.
 - Implemented Advanced RAG (LightRAG & NanoVectorDB) to maintain a persistent knowledge graph of historical “Smart Money” signals, improving signal-to-noise ratio in market anomaly detection.
 - Tech: Python 3.11, OpenAI/Gemini 2.0, LightRAG, NanoVectorDB, FastAPI, Streamlit, Apify, Git.
- ClimaChain** — AI-Based Weather Prediction System Mar 2026
- Engineered a Random Forest model in WEKA achieving ~85% accuracy and optimized F1-score performance; used Scikit-learn for class-imbalance handling and feature scaling.
 - Deployed a full-stack, responsive application on Vercel with PostgreSQL/Supabase for persistent data storage, including a frontend for user input and prediction visualization.

INTERNSHIP EXPERIENCE

- FlyRank.ai** — Machine Learning Intern (Remote) Jul–Sep 2026
- Built and validated a random-forest classifier to prioritize content pages for review from 30K real production pages across 32 clients, beating a rule-based baseline on every metric (F1 0.82 vs. 0.63; precision@50 = 1.00 on held-out data).
 - Designed a client-grouped train/test split and ran leakage audits to catch label reconstruction, cutting an initially inflated 0.87 accuracy to an honest, defensible 0.81 — caught and disclosed a false 0.9997 result during a sanity check.
 - Delivered a reason-coded, tiered priority queue (7 rule types, P1–P4 severity) turning model output into an actionable editor workflow, not just a score.

EDUCATION

- Bachelor of Engineering — Computer Engineering (IoT & Cybersecurity incl. Blockchain Technology)** — M.H. Saboo Siddik College of Engineering, Mumbai University 2024 – Expected May 2028
CGPA: 9.28/10 | Sem 4 SGPA: 9.7
- Senior Secondary (XII)** — Guru Nanak Khalsa College of Arts, Commerce & Science June 2024
Percentage: 80.00%

TRAININGS & CERTIFICATIONS

- Tata Group Forge — AI-powered data analytics job simulation at Tata iQ (January 2026).
- Aptech Learning — Data Analytics & Business Intelligence Analysis, Grade: Distinction, 88% weighted score (Jun 2025 – Apr 2026).